

#UpSkillWithKalpesh

Day 20

Data Science Unlocked

From Zero to Data Hero

SGD Classifier in Machine Learning



Kalpesh Pathade
@DataSimplified

SGD Classifier in Machine Learning

▼ Type

@DataSimplified

SGD Classifier in Machine Learning

The **SGD (Stochastic Gradient Descent) Classifier** is an efficient machine learning algorithm for **linear classification** that is particularly useful when dealing with **large datasets**. It is part of the **scikit-learn** library in Python.

1. What is Stochastic Gradient Descent (SGD)?

Gradient Descent (GD) is an optimization algorithm used to minimize the loss function in machine learning models by updating model parameters iteratively.

- **Batch Gradient Descent (BGD)**: Uses the entire dataset for each step of optimization.
- **Mini-batch Gradient Descent**: Uses small random subsets (mini-batches) of the dataset.
- **Stochastic Gradient Descent (SGD)**: Uses **only one data point** at a time for optimization.

💡 Why use SGD?

- Efficient for **large-scale datasets** (big data).
- Faster convergence in **online learning** scenarios.
- Works well with **sparse data** (like text classification).

2. Introduction to SGD Classifier

What is an SGD Classifier?

SGDClassifier is a **linear classifier** that optimizes **hinge loss** for SVM (Support Vector Machines) and **log loss** for logistic regression, among other loss functions.

It is available in **scikit-learn** as:

```
from sklearn.linear_model import SGDClassifier
```

Key Features of SGDClassifier

- Supports **multiple loss functions** (e.g., hinge, log, modified Huber).
- Uses **regularization** to prevent overfitting.
- Supports **different learning rate strategies**.

3. How SGD Classifier Works

Mathematical Understanding

SGDClassifier is based on the **linear model equation**:

$$y = wX + b$$

where:

- w = weight vector (learned during training)
- X = input features
- b = bias term

At each iteration, **SGD updates the weights** using the gradient of the loss function:

$$w = w - \eta \cdot \nabla L(w)$$

where:

- η = learning rate
- $\nabla L(w)$ = gradient of the loss function

Since SGDClassifier uses **stochastic** updates, it makes quick progress but can be noisy.

4. SGD Classifier in Python

4.1 Importing Required Libraries

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import accuracy_score, classification_report
```

4.2 Loading and Preparing the Data

```
# Generating synthetic data
from sklearn.datasets import make_classification
X, y = make_classification(n_samples=1000, n_features=20, random_state=42)

# Splitting the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Feature Scaling (important for SGD)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

4.3 Training the SGD Classifier

```
# Creating the SGD Classifier
sgd_clf = SGDClassifier(loss="hinge", penalty="l2", max_iter=1000, tol=1e-3, random_state=42)
```

```
# Training the model
sgd_clf.fit(X_train, y_train)
```

4.4 Making Predictions

```
# Predicting on test data
y_pred = sgd_clf.predict(X_test)

# Evaluating performance
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

5. Important Hyperparameters

5.1 Loss Functions

SGD supports multiple loss functions:

Loss Function	Description
"hinge"	SVM classifier (default)
"log"	Logistic Regression
"modified_huber"	Robust to outliers
"squared_hinge"	Variant of hinge loss

Example:

```
sgd_clf = SGDClassifier(loss="log") # Logistic Regression
```

5.2 Regularization (Penalty)

Regularization prevents overfitting by penalizing large coefficients.

Regularization	Description
"l2"	Ridge regression (default)
"l1"	Lasso regression (sparse weights)

"elasticnet"	Combination of L1 and L2
--------------	--------------------------

Example:

```
sgd_clf = SGDClassifier(penalty="elasticnet", l1_ratio=0.15)
```

5.3 Learning Rate Strategies

Learning Rate	Description
"constant"	Fixed learning rate
"optimal"	Predefined strategy
"invscaling"	Decreases as training progresses
"adaptive"	Adjusts automatically

Example:

```
sgd_clf = SGDClassifier(eta0=0.01, learning_rate="adaptive")
```

6. Advantages & Disadvantages

✓ Advantages

- ✓ Works well with **large datasets**.
- ✓ Efficient for **online learning**.
- ✓ Supports **various loss functions**.
- ✓ Handles **sparse and high-dimensional data**.

✗ Disadvantages

- ✗ Sensitive to **learning rate** selection.
- ✗ Convergence can be **noisy**.
- ✗ Requires **feature scaling** (Standardization).

7. Best Practices for Using SGD Classifier

- ✓ **Scale your data** using `StandardScaler()`.
- ✓ **Tune the learning rate** to ensure convergence.
- ✓ **Choose the right loss function** based on your task (e.g., `"log"` for classification).
- ✓ **Use regularization** (`l1`, `l2`, `elasticnet`) to prevent overfitting.
- ✓ **Increase** `max_iter` **and decrease** `tol` for better convergence.

8. When to Use SGD Classifier?

- When working with **large datasets** where standard classifiers (like SVM) are too slow.
- When performing **online learning** (learning as new data arrives).
- When handling **text classification** (e.g., spam detection, sentiment analysis).
- When dealing with **high-dimensional sparse data** (e.g., TF-IDF vectors).

9. Comparison with Other Models

Model	Pros	Cons
SGDClassifier	Fast, handles large data	Sensitive to learning rate
SVM (Support Vector Machine)	High accuracy, good for small datasets	Computationally expensive
Logistic Regression	Simple, interpretable	Not suitable for high-dimensional data

10. Summary

- **SGDClassifier** is a linear model optimized using **stochastic gradient descent**.
- It is best suited for **large datasets and online learning**.
- It supports **different loss functions**, including SVM and logistic regression.
- **Regularization and learning rate tuning** are crucial for optimal performance.
- **Feature scaling is necessary** before training an SGDClassifier.

11. Further Reading

📌 **Scikit-learn Documentation:** https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html

📌 **Gradient Descent Explanation:** https://ml-cheatsheet.readthedocs.io/en/latest/gradient_descent.html
